

Data Visualization

March 8, 2024

1 Milestone 1: data Visualization

This file contains some general data exploratory codes, comments and plots. We further used the visualizations in this file to make some planned decisions about feature engineering.

```
[38]: from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import LabelEncoder
      from sklearn.linear_model import LogisticRegression
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.svm import SVC
      from sklearn.model_selection import cross_val_score, StratifiedKFold
      from sklearn.model_selection import cross_val_predict
      from sklearn.metrics import roc_curve, roc_auc_score
      import pandas as pd
      import seaborn as sns
      import matplotlib.pyplot as plt
```

```
[39]: df = pd.read_csv("Assign1Data.csv")
      df.head()
```

```
[39]:   Gender Senior Citizen Partner  Tenure Months Phone Service Multiple Lines \
0    Male                No      No           2           Yes              No
1  Female                No      No           2           Yes              No
2  Female                No      No           8           Yes              Yes
3  Female                No     Yes          28           Yes              Yes
4    Male                No      No          49           Yes              Yes
```

```
   Internet Service Online Security Online Backup Device Protection \
0             DSL                Yes           Yes              No
1   Fiber optic                No           No              No
2   Fiber optic                No           No              Yes
3   Fiber optic                No           No              Yes
4   Fiber optic                No           Yes              Yes
```

```
   Tech Support Streaming TV Streaming Movies      Contract \
0         No          No           No  Month-to-month
1         No          No           No  Month-to-month
```

2	No	Yes	Yes	Month-to-month
3	Yes	Yes	Yes	Month-to-month
4	No	Yes	Yes	Month-to-month

	Paperless Billing	Payment Method	Monthly Charges \
0	Yes	Mailed check	53.85
1	Yes	Electronic check	70.70
2	Yes	Electronic check	99.65
3	Yes	Electronic check	104.80
4	Yes	Bank transfer (automatic)	103.70

	Total Charges	Churn Value
0	108.15	1
1	151.65	1
2	820.50	1
3	3046.05	1
4	5036.30	1

```
[40]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                7043 non-null  object
1   Senior Citizen        7043 non-null  object
2   Partner               7043 non-null  object
3   Tenure Months         7043 non-null  int64
4   Phone Service         7043 non-null  object
5   Multiple Lines        7043 non-null  object
6   Internet Service      7043 non-null  object
7   Online Security       7043 non-null  object
8   Online Backup         7043 non-null  object
9   Device Protection     7043 non-null  object
10  Tech Support          7043 non-null  object
11  Streaming TV          7043 non-null  object
12  Streaming Movies      7043 non-null  object
13  Contract              7043 non-null  object
14  Paperless Billing      7043 non-null  object
15  Payment Method        7043 non-null  object
16  Monthly Charges       7043 non-null  float64
17  Total Charges         7043 non-null  float64
18  Churn Value           7043 non-null  int64
dtypes: float64(2), int64(2), object(15)
memory usage: 1.0+ MB
```

```
[41]: print(df.columns)
```

```
Index(['Gender', 'Senior Citizen', 'Partner', 'Tenure Months', 'Phone Service',
      'Multiple Lines', 'Internet Service', 'Online Security',
      'Online Backup', 'Device Protection', 'Tech Support', 'Streaming TV',
      'Streaming Movies', 'Contract', 'Paperless Billing', 'Payment Method',
      'Monthly Charges', 'Total Charges', 'Churn Value'],
      dtype='object')
```

Here we printed all the categories there are in each columns.

```
[42]: for col in df.columns:
      print(f"Column {col}, categories {df[col].unique()}")
```

```
Column Gender, categories ['Male' 'Female']
Column Senior Citizen, categories ['No' 'Yes']
Column Partner, categories ['No' 'Yes']
Column Tenure Months, categories [ 2  8 28 49 10  1 47 17  5 34 11 15 18  9  7
12 25 68 55 37  3 27 20  4
58 53 13  6 19 59 16 52 24 32 38 54 43 63 21 69 22 61 60 48 40 23 39 35
56 65 33 30 45 46 62 70 50 44 71 26 14 41 66 64 29 42 67 51 31 57 36 72
0]
Column Phone Service, categories ['Yes' 'No']
Column Multiple Lines, categories ['No' 'Yes' 'No phone service']
Column Internet Service, categories ['DSL' 'Fiber optic' 'No']
Column Online Security, categories ['Yes' 'No' 'No internet service']
Column Online Backup, categories ['Yes' 'No' 'No internet service']
Column Device Protection, categories ['No' 'Yes' 'No internet service']
Column Tech Support, categories ['No' 'Yes' 'No internet service']
Column Streaming TV, categories ['No' 'Yes' 'No internet service']
Column Streaming Movies, categories ['No' 'Yes' 'No internet service']
Column Contract, categories ['Month-to-month' 'Two year' 'One year']
Column Paperless Billing, categories ['Yes' 'No']
Column Payment Method, categories ['Mailed check' 'Electronic check' 'Bank
transfer (automatic)'
'Credit card (automatic)']
Column Monthly Charges, categories [ 53.85  70.7   99.65 ... 108.35  63.1   78.7
]
Column Total Charges, categories [ 108.15  151.65  820.5   ... 7362.9   346.45
6844.5 ]
Column Churn Value, categories [1 0]
```

Here the describe function only showed the attributes of the columns which had numerical data. Later we converted all the categorical values in integers and used describe again. That can be seen further down.

```
[43]: df.describe()
```

```
[43]:
```


	Tenure Months	Monthly Charges	Total Charges	Churn Value
count	7043.000000	7043.000000	7043.000000	7043.000000
mean	32.371149	64.761692	2283.300441	0.265370

std	24.559481	30.090047	2265.000258	0.441561
min	0.000000	18.250000	18.800000	0.000000
25%	9.000000	35.500000	402.225000	0.000000
50%	29.000000	70.350000	1400.550000	0.000000
75%	55.000000	89.850000	3786.600000	1.000000
max	72.000000	118.750000	8684.800000	1.000000

```
[44]: # Showing Churn distribution

sns.set(style="whitegrid")

sns.countplot(x='Churn Value', data=df)
plt.title('Churn Distribution')
plt.show()
```



This bar chart illustrates the distribution of churn. The significant imbalance between the retained and churned customer is crucial to consider when selecting performance metrics for model evaluation. Models like Random Forest Classifier and Decision Tree are insensitive to such imbalances. We can use techniques like SMOTE for data balancing before training more sensitive models such as Logistic Regression or SVM.

```

[45]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

plt.figure(figsize=(20, 17))

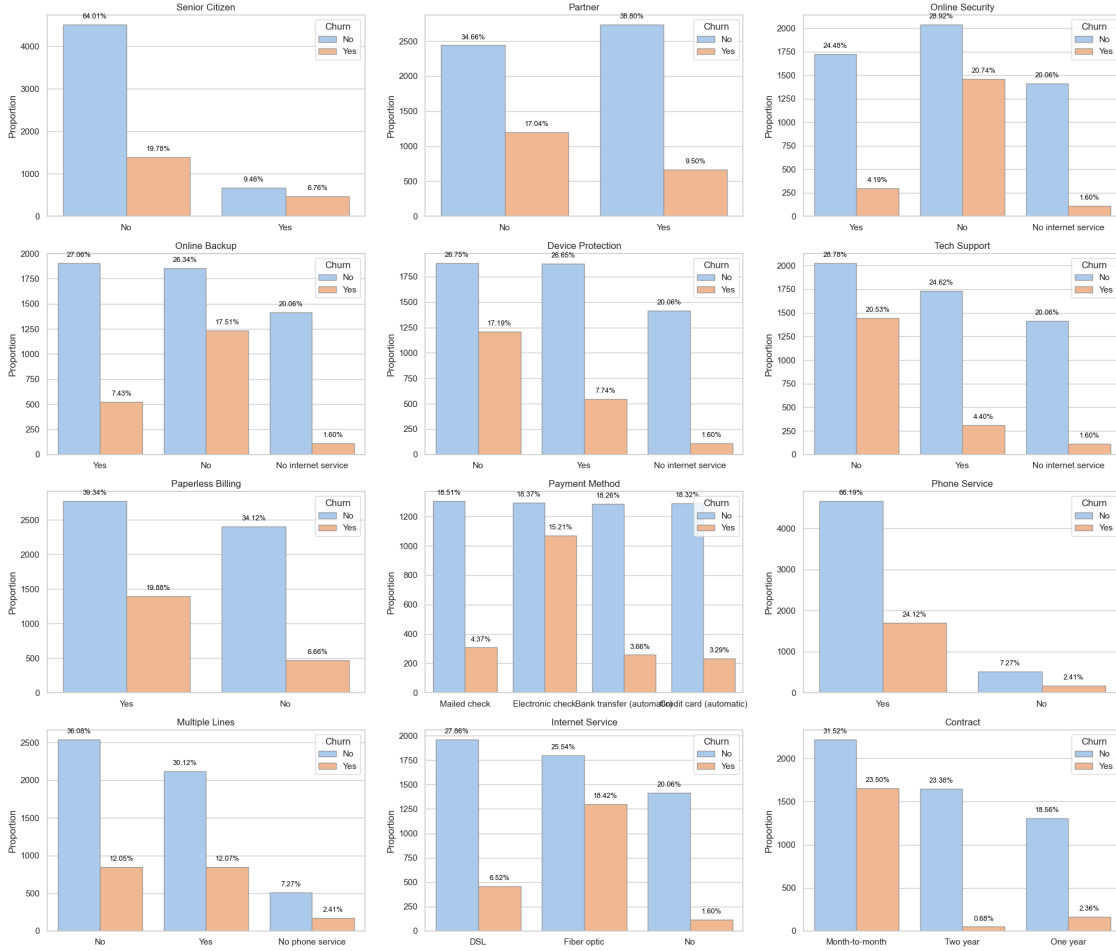
columns_to_plot = [
    'Senior Citizen', 'Partner', 'Online Security',
    'Online Backup', 'Device Protection', 'Tech Support',
    'Paperless Billing', 'Payment Method',
    'Phone Service', 'Multiple Lines', 'Internet Service', 'Contract'
]

for i, column in enumerate(columns_to_plot, 1):
    plt.subplot(4, 3, i) #grid dimension
    ax = sns.countplot(x=column, hue='Churn Value', data=df, palette='pastel', □
    ↪edgecolor='.6')
    plt.title(column.replace('_', ' ').title())
    plt.xlabel('')
    plt.ylabel('Proportion')
    plt.legend(title='Churn', labels=['No', 'Yes'], loc='upper right')
    total = len(df[column])

    # adding the normalized percentage on top of the bars
    for p in ax.patches:
        height = p.get_height()
        percentage = '{:1.2f}%'.format(100 * height/total)
        ax.annotate(percentage, (p.get_x() + p.get_width() / 2., height), □
        ↪ha='center', va='center', fontsize=9, color='black', xytext=(0, 10), □
        ↪textcoords='offset points')

plt.tight_layout()
plt.show()

```



The churn rate of customers is influenced by various factors, including age, gender, and payment method. Almost half of the senior citizen can be seen to be churned, suggesting the need for more effort to retain senior citizen by providing more comfort. Customers with partners are slightly less likely to churn, suggesting services that appeal to couples might keep them loyal. Online security services significantly decrease churn, highlighting the importance of security features. Tech support also significantly reduces churn, underlining the importance of customer service. Payment methods also impact churn rates, especially electronic check, they seem to be causing comparatively large effect on customer churn. The set of plots above give tons of such insights, to make it more easily readable we plotted a churn correlation plot in the end which shows how much a particular column as a whole effects churn.

```
[46]: import matplotlib.pyplot as plt
import seaborn as sns

service_columns = ['Multiple Lines', 'Online Security',
                  'Online Backup', 'Device Protection', 'Tech Support',
                  'Streaming TV', 'Streaming Movies']
```

```

plt.figure(figsize=(10, 14))

for i, column in enumerate(service_columns, 1):
    ax = plt.subplot(len(service_columns), 1, i)
    churn_rate = df.groupby(column)['Churn Value'].value_counts(normalize=True).
    ↪unstack()
    churn_rate.plot(kind='bar', stacked=True, ax=ax)
    plt.title(column)
    plt.xlabel('Service')
    plt.ylabel('Churn Rate')
    plt.legend(title='Churn', bbox_to_anchor=(1, 1))


plt.tight_layout()
plt.show()

```

```

/var/folders/4z/qh15q30n25nckhc10gfcfdwh0000gn/T/ipykernel_30291/4077526128.py:1
9: UserWarning: Tight layout not applied. tight_layout cannot make axes height
small enough to accommodate all axes decorations.
    plt.tight_layout()

```



The plot shows above have some of its label hidden, The ones in the middle are “No phone service”

and “No Internet service”. The columns which are plotted had 3 categories in them, multiple lines had “No phone service” value in it if the Phone service column had “No” in that row. As for the rest of the services they had a categorical value “No internet service” in the row where the Internet Service column had the value “No”. Here we just plotted all those categories and the effect they had on churn to decide if any can be further worked upon during feature engineering to optimize our models performance.

```
[47]: # Converting the categorical values into integers

for col in df.select_dtypes(include=['object']).columns:
    df[col] = df[col].astype('category').cat.codes

df.head()
```

```
[47]:
```

	Gender	Senior Citizen	Partner	Tenure	Months	Phone Service	\
0	1	0	0		2	1	
1	0	0	0		2	1	
2	0	0	0		8	1	
3	0	0	1		28	1	
4	1	0	0		49	1	

	Multiple Lines	Internet Service	Online Security	Online Backup	\
0	0	0	2	2	
1	0	1	0	0	
2	2	1	0	0	
3	2	1	0	0	
4	2	1	0	2	

	Device Protection	Tech Support	Streaming TV	Streaming Movies	Contract	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	2	0	2	2	0	
3	2	2	2	2	0	
4	2	0	2	2	0	

	Paperless Billing	Payment Method	Monthly Charges	Total Charges	\
0	1	3	53.85	108.15	
1	1	2	70.70	151.65	
2	1	2	99.65	820.50	
3	1	2	104.80	3046.05	
4	1	0	103.70	5036.30	

	Churn Value
0	1
1	1
2	1
3	1

```
[49]: df.describe()
```

```
[49]:
```

	Gender	Senior Citizen	Partner	Tenure Months	Phone Service \
count	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000
mean	0.504756	0.162147	0.483033	32.371149	0.903166
std	0.500013	0.368612	0.499748	24.559481	0.295752
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	9.000000	1.000000
50%	1.000000	0.000000	0.000000	29.000000	1.000000
75%	1.000000	0.000000	1.000000	55.000000	1.000000
max	1.000000	1.000000	1.000000	72.000000	1.000000

	Multiple Lines	Internet Service	Online Security	Online Backup \
count	7043.000000	7043.000000	7043.000000	7043.000000
mean	0.940508	0.872923	0.790004	0.906432
std	0.948554	0.737796	0.859848	0.880162
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	1.000000	1.000000	1.000000	1.000000
75%	2.000000	1.000000	2.000000	2.000000
max	2.000000	2.000000	2.000000	2.000000

	Device Protection	Tech Support	Streaming TV	Streaming Movies \
count	7043.000000	7043.000000	7043.000000	7043.000000
mean	0.904444	0.797104	0.985376	0.992475
std	0.879949	0.861551	0.885002	0.885091
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	1.000000	1.000000	1.000000	1.000000
75%	2.000000	2.000000	2.000000	2.000000
max	2.000000	2.000000	2.000000	2.000000

	Contract	Paperless Billing	Payment Method	Monthly Charges \
count	7043.000000	7043.000000	7043.000000	7043.000000
mean	0.690473	0.592219	1.574329	64.761692
std	0.833755	0.491457	1.068104	30.090047
min	0.000000	0.000000	0.000000	18.250000
25%	0.000000	0.000000	1.000000	35.500000
50%	0.000000	1.000000	2.000000	70.350000
75%	1.000000	1.000000	2.000000	89.850000
max	2.000000	1.000000	3.000000	118.750000

	Total Charges	Churn Value
count	7043.000000	7043.000000
mean	2283.300441	0.265370

std	2265.000258	0.441561
min	18.800000	0.000000
25%	402.225000	0.000000
50%	1400.550000	0.000000
75%	3786.600000	1.000000
max	8684.800000	1.000000

As promised, describing the data frame after converting the categorical values into integers. We can see mean to find out if the column is balanced. Min and max to see how many categories are there in each columns. Using the percentiles we can also have an idea as to how the data is distributed.

```
[48]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestClassifier

X = pd.get_dummies(df.drop('Churn Value', axis=1))
y = df['Churn Value']

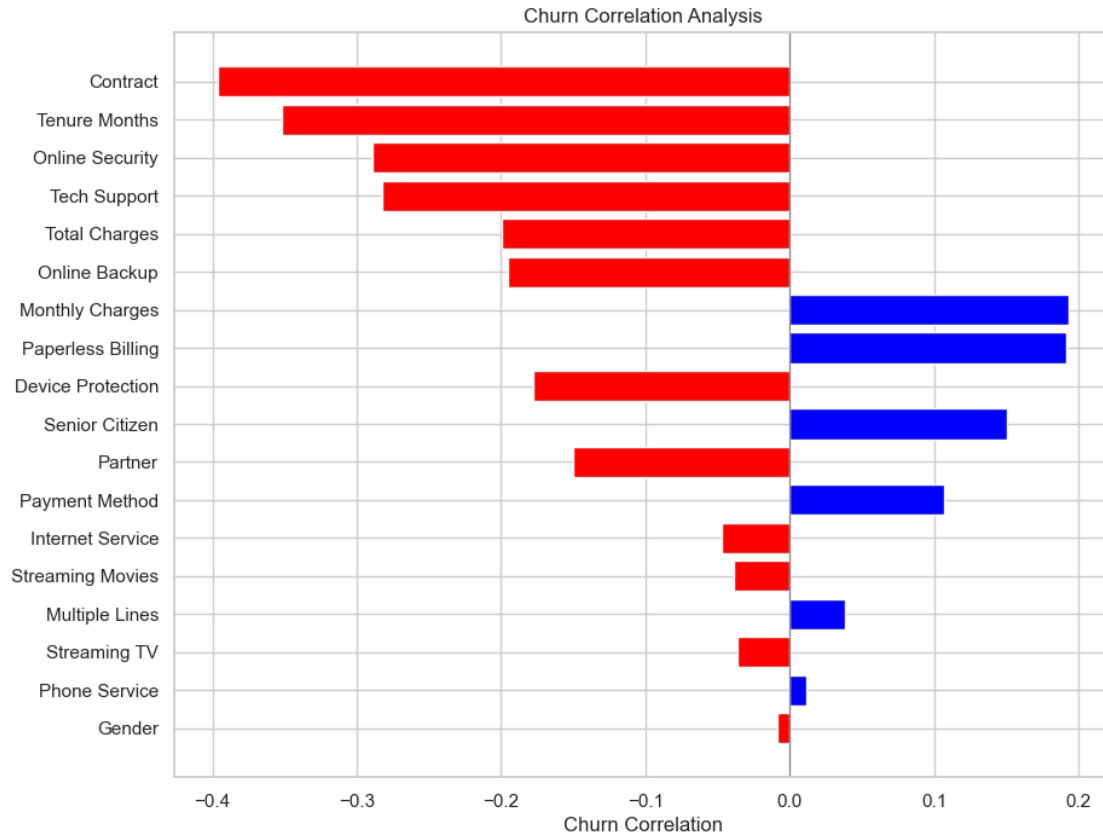
# Calculate the correlation of each feature with the target
correlations = X.apply(lambda x: x.corr(df['Churn Value']))

# create a dataframe of the correlations calculated using the lambda function
# above
features_df = pd.DataFrame({'Correlation': correlations})

# Sort the DataFrame based on the absolute values of the correlations
features_df = features_df.sort_values(by='Correlation', key=abs, ascending=True)

# Create the color map based on the correlation values
colors = ['red' if x < 0 else 'blue' for x in features_df['Correlation']]

# Plot the figure
plt.figure(figsize=(10, 8))
bars = plt.barh(features_df.index, features_df['Correlation'], color=colors)
plt.xlabel('Churn Correlation')
plt.title('Churn Correlation Analysis')
plt.axvline(x=0, color='grey', linewidth=0.8)
plt.show()
```



The horizontal bar chart presented, shows the correlation of various features with Churn Value. Features with positive correlations are colored blue, they suggest a positive correlation and a tendency to increase churn possibility (simply leading to customer leaving Telco company). On the other hand, features with negative correlations are colored red and are suggesting a potential retention factors (it can be noted that most of the columns are leading to retention and thus our data set is imbalanced). The graph is sorted to prioritize features with stronger correlations on the top and weaker on the bottom. This helps in feature selection and provides insights into areas that may require more focused customer retention strategies to further decrease customer churn.